

Understanding Einsum

And it's decomposition in terms of primitive operations

Nikita Vedeneev

nikitaved@github

Preliminaries

- **Index set** $I(d) := \{0, \dots, d - 1 | d > 0\}$
- Two index sets $I(a)$ and $J(b)$ are **equal**, i.e. $I(a) = J(b)$ if $a = b \vee a = 1 \vee b = 1$, or, alternatively, when the shapes $(a,)$ and $(b,)$ **broadcast**

Sets of index sets

Consider $I(1)$ and $J(3)$. Since $(1,)$ and $(3,)$ broadcast, $I(1) = J(3)$, so $\{I(1), J(3)\} = \{I(1)\} = \{J(3)\}$.

Indexing arrays

- Suppose a is a 2D array of shape (s_1, s_2) , and $I(s_1), J(s_2)$ are index sets, then

$$a[I, J] = a[I(s_1), J(s_2)] := \text{a sequence } (a[i, j])_{ij} \forall (i, j) \in I(s_1) \times J(s_2),$$

where \times denotes the Cartesian product.

Note that parameters s_1, s_2 are made implicit for lighter notation.

- If $I = J$, (i.e. $I(s_1) = J(s_2)$), then

$$a[I, I] := \left(\begin{matrix} a[i, j], i = j, \\ 0, i \neq j. \end{matrix} \right)_{ij}, \forall (i, j) \in i \times j$$

Note on indexing with equal index sets

If $I(s_1) = J(s_2)$, then $I(s_1) \times J(s_2) = J(s_2) \times I(s_1)$, and this implies $s_1 = s_2$.
This means we cannot use same index sets to index dimensions with different lengths!

- Of course, we can generalize indexing with index sets to more than just 2 dimensions

$$\begin{aligned} a[I, I, J, J] &= (a[i_1, i_2, j_1, j_2])_{i_1 i_2 j_1 j_2} \\ &= \left(\begin{cases} a[i_1, i_2, j_1, j_2], i_1 = i_2, j_1 = j_2, \\ 0, \text{ otherwise.} \end{cases} \right)_{i_1 i_2 j_1 j_2} \\ &\quad \forall (i_1, i_2, j_1, j_2) \in I \times I \times J \times J. \end{aligned}$$

with implicit constraints

`a.shape[0] = a.shape[1]` and `a.shape[2] = a.shape[3]`

Labelling

Suppose x is a n -dim array

- A tuple $(L_1, \dots, L_n) \in \mathbb{R}^n$ is a **labelling** of dimensions in x
- So we can replace **index sets** with **labels** when indexing, i.e.
 $x[L_1, \dots, L_n] := x[l_1, \dots, l_n]$, where $L_i = L_j \implies l_i = l_j$
- Extension to **binary** ops

$$x[L_1^x, \dots, L_i^x] \cdot y[L_1^y, \dots, L_j^y] := x[l_1^x, \dots, l_i^x] \cdot y[l_1^y, \dots, l_j^y]$$

$$\text{where } L_k^x = L_l^y \implies l(x.shape[k]) = l(y.shape[l])$$

$$\wedge [k = l \vee k = 1 \vee l = 1 : \forall (k, l) \in l(x.shape[k]) \times l(y.shape[l])]$$

Sequences and arrays

$x[l]$, $y[l]$ are sequences of real numbers, so $x[l] \cdot y[l]$ is well-defined and is a sequence. Moreover, there is a **bijection** between sequences and arrays (they are **isomorphic**), so we can use these objects interchangeably.

Suppose x is a n -dim array

- A **contraction** of x over a **contraction set** (of labels) C with a **saving set** (of labels) S is defined as

$$\sum_{C \setminus S} x[L_1, \dots, L_n] := \left(\sum_{i_k, k \in \{k | L_k \in C\} \setminus \{\max\{k | L_k \in S\}\}} x[i_1, \dots, i_n] \right)_{(i_k, k \in \{k | L_k \notin C\} \cup \{\max\{k | L_k \in S\}\})}$$
$$\forall (i_1, \dots, i_n) \in I_1 \times \dots \times I_n$$

Einsum. Decomposition

Given

- n arrays x_1, \dots, x_n of ranks r_1, \dots, r_n
- and their corresponding labellings $(L_1^{x_1}, \dots, L_{r_1}^{x_1}), \dots, (L_1^{x_n}, \dots, L_{r_n}^{x_n})$
- and the output's labelling (L_1^o, \dots, L_k^o) such that
 - $|\{L_1^o, \dots, L_k^o\}| = k$
 - $\forall L_i^o \exists L_k^{x_j} : L_i^o = L_k^{x_j}$

`einsum` computes

$$o[\pi(L_1^o, \dots, L_k^o)] = \sum_{\cup_i \{L_1^{x_i}, \dots, L_{r_i}^{x_i}\} \setminus \{L_1^o, \dots, L_k^o\}} \prod_{j=1}^n x_j[L_1^{x_j}, \dots, L_{r_j}^{x_j}]$$

for some permutation of labels π .

Einsum is powerful

- $x \in \mathbb{R}^{n \times n}, C = \{I\}, S = \emptyset \implies \sum_{C \setminus S} x[I, I] \equiv \sum_{ii} x[i, i] = \text{trace}(x)$
- $x \in \mathbb{R}^{n \times n}, C = \{I\}, S = \{I\} \implies \sum_{C \setminus S} x[I, I] \equiv \text{diag}(x)$
- $x \in \mathbb{R}^{m \times n}, y \in \mathbb{R}^{n \times p}, C = \{N\}, S = \emptyset \implies \sum_{C \setminus S} x[M, N] \cdot y[N, P] \equiv \text{matmul}(x, y)$
- $x \in \mathbb{R}^m, y \in \mathbb{R}^n, C = \emptyset, S = \emptyset \implies \sum_{C \setminus S} x[M] \cdot y[N] = (x[i] \cdot y[j])_{ij} \equiv \text{outer}(x, y)$

Einsum. Decomposition

`einsum` generalization to multiple inputs is possible because

- \sum and \prod **distribute**
- \prod/\cdot is a **binary** operation that is also **associative**

As such, `einsum` is a composition of

- "unary" contractions $\sum_{C \setminus S} x[L_1, \dots, L_n]$
- "binary" contractions $\sum_{C \setminus S} x[L_1^x, \dots, L_m^x] \cdot y[L_1^y, \dots, L_n^y]$

Associativity - contraction order matters

Consider a sequence of matrix multiplications of shapes $(1, n), (n, n^2), (n^2, n^3)$. Multiplying from left to right will yield $O(n^5)$ and from right to left - $O(n^6)$. We use the `opt_einsum` package to generate better than naive orderings.

"binary" contractions $\sum_{C \setminus S} x[L_1^x, \dots, L_m^x] \cdot y[L_1^y, \dots, L_n^y]$ are most interesting.

- Let $L_x = \{L_1^x, \dots, L_m^x\}, L_y = \{L_1^y, \dots, L_n^y\}$
- $U_x = L_x \cap C \setminus L_y$ and $U_y = L_y \cap C \setminus L_x$ - **unique** to x and y labels

U_x and U_y can be "unary" contracted on the spot. So let's assume $U_x = U_y = \emptyset$.

"binary" contractions $\sum_{C \setminus S} x[L_1^x, \dots, L_m^x] \cdot y[L_1^y, \dots, L_n^y]$ are most interesting. Output labels are either provided, or deduced, so $L_o = \{L_1^o, \dots, L_p^o\}$ is defined.

- Let $L_b = L_o \cap L_x \cap L_y$ the **sorted** set of **batch** labels
- Let $L_C^x \subseteq L_x \cap C$ the **sorted** set of contraction labels in x where the order is determined by the dimension the labels appear in.
Same for L_C^y
- Let $L_R^x \subseteq L_x \setminus (L_b \cup L_C^x)$ the **sorted** set of the remaining labels with the same order structure as L_C^x . Same for L_R^y

"binary" contractions $\sum_{C \setminus S} x[L_1^x, \dots, L_m^x] \cdot y[L_1^y, \dots, L_n^y]$ are most interesting

- **partition/align** labelled dimensions such that

$$x[L_1^x, \dots, L_m^x] \mapsto \tilde{x}[L_b, L_R^x, L_C^x]$$

$$y[L_1^y, \dots, L_n^y] \mapsto \tilde{x}[L_b, L_C^y, L_R^y]$$

- **flatten** dimensions $L_R^x, L_C^x, L_C^y, L_R^y$, do matmul, unflatten back to get an array with the labelling (L_b, L_R^x, L_R^y)
- (optionally) permute dimensions in (L_b, L_R^x, L_R^y) to match output

Einsum. Decomposition. Contraction with reductions

"binary" contractions $\sum_{C \setminus S} x[L_1^x, \dots, L_m^x] \cdot y[L_1^y, \dots, L_n^y]$ are most interesting

- **partition/align** labelled dimensions such that

$$x[L_1^x, \dots, L_m^x] \mapsto \tilde{x}[L_b, L_R^x, L_C^x]$$

$$y[L_1^y, \dots, L_n^y] \mapsto \tilde{y}[L_b, L_R^y, L_C^y]$$

- **reshape** \tilde{x}, \tilde{y}

$$(L_b, L_R^x, L_C^x) \mapsto (L_b, L_R^x, \underbrace{1, \dots, 1}_{|L_R^y| \text{ times}}, L_C^x)$$

$$(L_b, L_R^y, L_C^y) \mapsto (L_b, \underbrace{1, \dots, 1}_{|L_R^x| \text{ times}}, L_R^y, L_C^y)$$

- do elem-wise **mul** and **sum-out** last $|C|$ dims to get an array with the labelling (L_b, L_R^x, L_R^y)
- (optionally) permute dimensions in (L_b, L_R^x, L_R^y) to match output